Evaluation Report

E-Commerce Return Fraud Detection

Dataset: E-Commerce Return Fraud Detection

Records: 50,500 raw → 50,000 after cleaning

Features: 30 raw → 41 final modeling features

Target: return_fraud (0 = Genuine, 1 = Fraud)

Documented By: Ragavendra R

AI/ML Engineer

EVALUATION REPORT

1. Class Imbalance Analysis

- Fraud: 35,132 (70.3%) | Genuine: 14,868 (29.7%) → 2.4:1 ratio.
- A naive classifier that always predicts fraud scores 70.3% accuracy without learning anything.
- Raw accuracy is a misleading metric — F1, ROC-AUC, PR-AUC, and Confusion Matrix used as primary metrics.
- Handling strategy:
  - Logistic Regression: class_weight='balanced' — auto-adjusts loss weights.
  - XGBoost: scale_pos_weight = genuine_count / fraud_count = 0.423 — equivalent to class weighting.
  - SMOTE tested as alternative to scale_pos_weight (see Section 3).

2. Model Comparison

Three models evaluated on identical 80/20 stratified train-test splits:

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
Logistic Regression	0.9137	0.9701	0.9051	0.9365	0.9754	0.9863
Random Forest	0.9995	0.9993	1.0000	0.9996	0.9999	0.9999
XGBoost (Selected)	0.9995	0.9994	0.9999	0.9996	1.0000	1.0000

- LR recall drops to 90.5% — the fraud boundary is non-linear (return-rate threshold); a linear model cannot learn it.
- RF and XGBoost are near-identical — XGBoost edges ahead at the 4th decimal consistently across all metrics.
- The 8.5-point accuracy gap between LR and tree models understates the real damage — the confusion matrix breakdown clarifies actual cost.

3. Imbalance Handling — scale_pos_weight vs SMOTE

Both methods applied to XGBoost only; test set always the original unmodified split:

Method	Accuracy	Precision	Recall	F1	FP	FN
XGB + scale_pos_weight	0.9995	0.9994	0.9999	0.9996	4	1
XGB + SMOTE	0.9994	0.9994	0.9997	0.9996	4	2

- Decision: scale_pos_weight selected.
- Both methods produce identical FP counts (4 false alarms on 2,974 genuine customers).
- scale_pos_weight: 1 false negative vs SMOTE: 2 false negatives — one fewer missed fraud.
- scale_pos_weight requires no data augmentation; SMOTE synthesizes 16,212 artificial samples introducing noise risk without measurable gain.

4. Confusion Matrix Analysis

Results on 10,000 test samples (2,974 genuine, 7,026 fraud):

Model	True Positive	False Positive	False Negative	True Negative
Logistic Regression	6,359	196	667	2,778
Random Forest	7,026	5	0	2,969
XGBoost	7,025	4	1	2,970

- LR: 667 false negatives — 667 fraudulent orders slip through unchecked.
- LR: 196 false positives — 196 genuine customers face unnecessary friction.
- RF: 0 false negatives (zero missed frauds) but 5 false positives.
- XGBoost: 1 FN and 4 FP — better probability calibration at the margin vs RF.

5. Business Cost Analysis

Assumptions: FN cost = ₹5,000 (unrecovered fraud loss) | FP cost = ₹1,000 (customer friction / investigation).

Model	FP	FN	Total Cost (₹)
Logistic Regression	196	667	3,531,000
Random Forest	5	0	5,000
XGBoost	4	1	9,000

- LR's ₹3.53M cost is 99.7% driven by false negatives — 667 missed frauds × ₹5,000 = ₹3.335M.
- XGBoost at ₹9,000 costs 1.8× more than RF but delivers better probability calibration for threshold tuning.
- Cost gap between LR and tree models: 390× on the same 10,000-row test set.

6. ROC and Precision-Recall Curve Observations

ROC Curve:

- XGBoost AUC = 1.000 | RF AUC = 0.9999 | LR AUC = 0.9754.
- XGBoost and RF both reach the top-left corner — near-perfect TPR across every FPR threshold.
- LR's curve bends away at high sensitivity — full recall with LR requires accepting a higher false positive rate.

Precision-Recall Curve:

- XGBoost AP = 1.000 | RF AP = 0.9999 | LR AP = 0.9863.
- All three models clear the 0.703 baseline (fraud prevalence) by a large margin.
- XGBoost and RF maintain precision near 1.0 all the way to maximum recall — threshold tuning is largely unnecessary.
- LR's curve drops off past ~95% recall — chasing full coverage requires accepting more false alarms.

7. SHAP Feature Importance (XGBoost)

Rank	Feature	Direction
1	return_rate	High rate → Fraud
2	ip_risk_score	High score → Fraud (context-dependent)
3	previous_returns	High count → Fraud
4	high_return_rate	Present → Fraud
5	support_intensity	High ratio → Fraud

- return_rate dominates by a wide margin — raw behavioral history outperforms all other signals.
- ip_risk_score ranks second but with wider SHAP spread — context-dependent, not a standalone predictor.
- Demographic features (gender, device_type, city_tier) cluster near zero — model ignored them entirely, confirming EDA findings.
- customer_age SHAP ≈ 0.00 — consistent with its 0.00 Pearson correlation to return_fraud.

8. Final Model Selection — XGBoost

- Highest ROC-AUC (1.000) and PR-AUC (1.000) among all three models.
- Fewer false positives than RF (4 vs 5) — fewer genuine customers incorrectly flagged.
- scale_pos_weight handles class imbalance without synthetic data generation.
- Gradient boosting learns non-linear feature interactions (return_rate × ip_risk_score) that Logistic Regression cannot.
- Natively supports SHAP explainability via TreeExplainer — critical for fraud operations team trust.
- 5-fold CV F1 = 0.9997 ± 0.0002 — confirms model is not overfitting to a single test split.

Note on Near-Perfect Scores:

- Scores near 1.0 are expected for synthetic data where fraud follows clean, rule-based patterns.
- In production with real data, noise and concept drift will lower all metrics — this is expected.
- Project value lies in the correct pipeline, leakage-free features, and reproducible workflow — not the inflated accuracy number.

9. Risk Categorization

Post-prediction, fraud probability is translated into actionable operational tiers:

Probability Range	Risk Category	Recommended Action
< 0.30	Low Risk	Auto-approve. Process refund without manual review.
0.30 – 0.70	Medium Risk	Manual review. Verify order history and return reason.
> 0.70	High Risk	Hold refund. Flag for fraud team. Check IP and return history.